

# Feature engineering in NLP

Mariana Romanyshyn  
Computational Linguist at *Grammarly*

# Contents

---

1. Presentation of text in NLP
2. Feature engineering
3. Feature encoding
4. A few words about ngrams (optional)
5. Logistic regression (optional)
6. Error correction use case

# 1. Presentation of text in NLP

# Presentation of text in NLP

---

*Trump beat Clinton in the election.*

**= or ≠**

*Clinton beat Trump in the election.*

# Presentation of text in NLP

- 
- Bag of words



# Presentation of text in NLP

- 
- Bag of words



- Sequence

I shot an elephant in my pajamas.

# Presentation of text in NLP

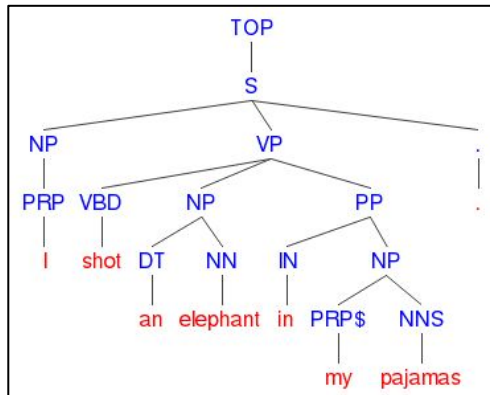
- Bag of words



- Sequence

I shot an elephant in my pajamas.

- Tree



# Presentation of text in NLP

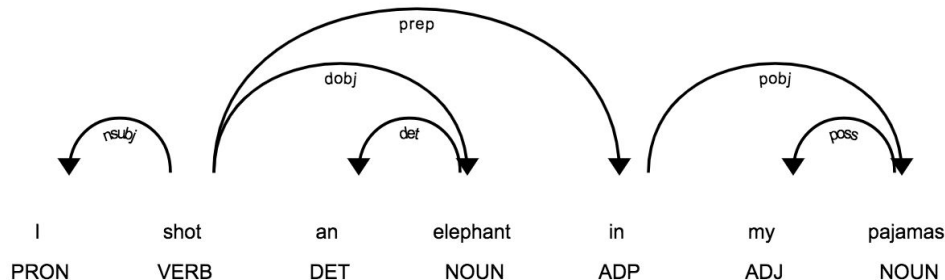
- Bag of words



- Sequence

I shot an elephant in my pajamas.

- Tree





# Presentation of text in NLP


- Bag of words



- Sequence

I shot an elephant in my pajamas.

- Tree



- Graph



# Presentation of text in NLP

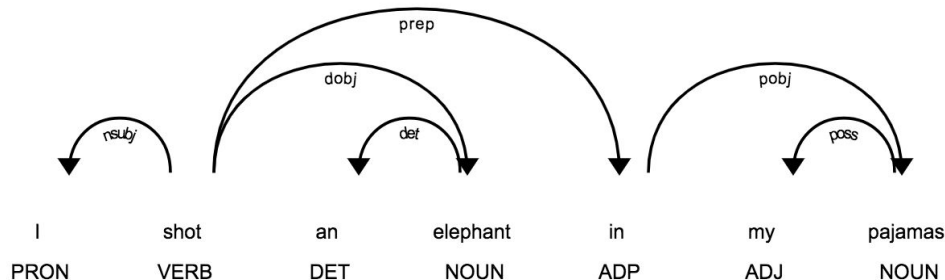
- Bag of words



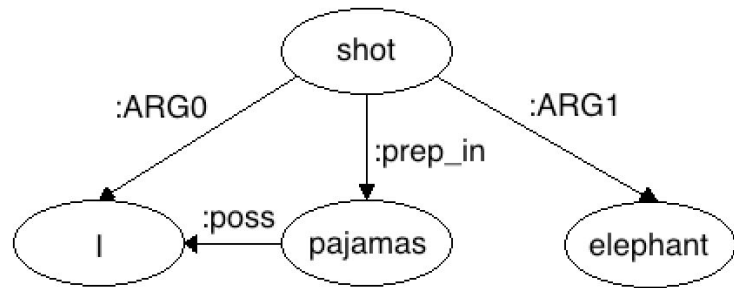
- Sequence

I shot an elephant in my pajamas.

- Tree



- Graph



# Think of different NLP tasks

---

How would you view the text if you need...

- classification of news articles by topic?
- named-entity recognition?
- sentiment assignment to objects in the text?
- abstractive text summarization?

## 2. Feature engineering

\_\_\_\_\_



# The Form

— — —

The form — how the word is written.

# Form Features


— — —

- Capitalization, hyphenation, apostrophes
- Lemma or stem
- Number of stems
- Number and types of affixes
- Length of the word/lemma/stem
- Number of tokens, position of a token
- Number of syllables in a word
- Ratio of vowels vs. consonants
- Voiced vs. voiceless consonants
- All possible frequencies

# Form Features

---

*a man-eating shark* vs. *a man eating shark*

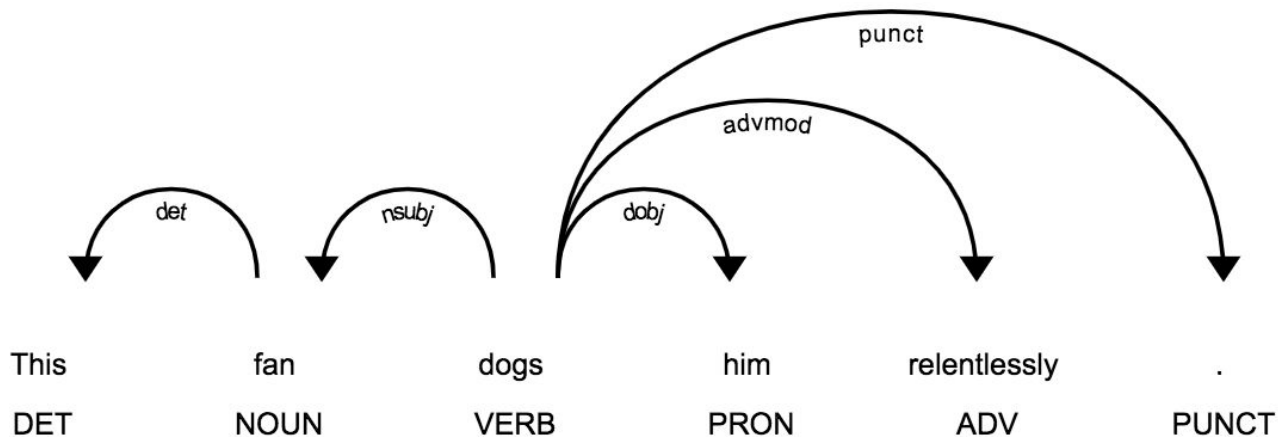




# The Function

---

The function — what the word does and how it interacts with other words in the text.



# Function Features

- — —
- Part of speech, part-of-speech tag
- Morphological properties:
  - *gender, animacy, number, person, case*
  - *aspect, voice, tense, degree of comparison*
- Constituents
  - *parents, children*
- Direct and indirect dependencies
  - *parents, children, type of relation*
- Depth of the syntactic tree
- Statistics: *POS+word, POS ngrams, syntactic ngrams*

# The Meaning

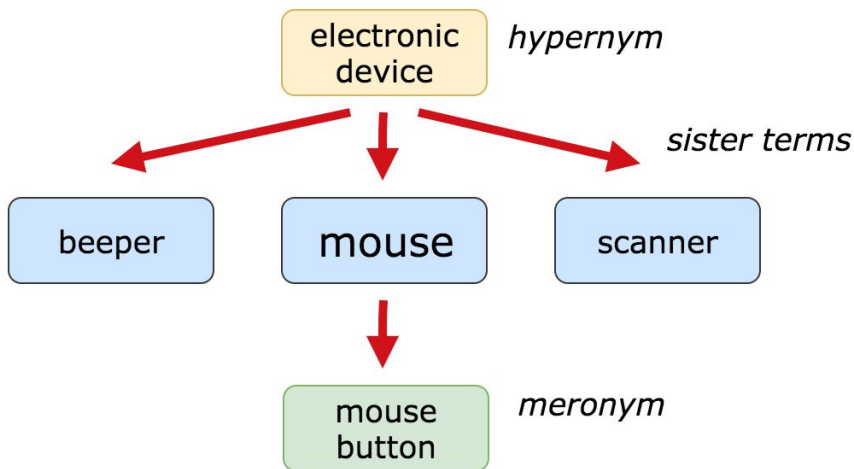

---

The meaning — the particular sense the word obtains in the context and how it correlates with other words and senses.



# Lexical Semantics in WordNet

— — —



# Entailment

— — —

Entailment examples:

- Mom slices a cucumber. => A woman cuts a vegetable.
- A jogger was spotted. => Someone was jogging.
- Her husband was snoring. => Her husband was sleeping.
- The king of France is bald. => There exists a king of France.

# Meaning Features

— — —

- Word sense
- Number of senses
- Shortest path to another word sense
- Similarity to another sense
- Synonyms/antonyms, hyponyms/hypernyms, meronyms/holonyms
- Entailment
- Semantic role

# 3. Feature encoding

# Encode as a bag of words: boolean

— — —

The pound extended losses against both the dollar and the euro .

[	1,	1,	0,	0,	1,	0,	...]
	the	dollar	hello	pirate	losses	run	



# Encode as a bag of words: count

— — —

The pound extended losses against both the dollar and the euro .

[ 3,	1,	0,	0,	1,	0,	...]
the	dollar	hello	pirate	losses	run	

# Encode as a bag of words: tf-idf

---

The pound extended losses against both the dollar and the euro .

[	0.1,	0.7,	0,	0,	0.4,	0,	...]
	the	dollar	hello	pirate	losses	run	

# Encode as a bag of features

— — —

DT NN VBD NNS IN DT DT NN CC DT NN .  
The pound extended losses against both the dollar and the euro .

“losses”/NNS:

{“word-2”: “pound”,  
“word-1”:  
“extended”,  
“word+1”: “against”,  
“word+2”: “both”,

“tag-2”: “NN”,  
“tag-1”: “VBD”,  
“tag+1”: “IN”,  
“tag+2”: “DT”}

# Encode as a bag of features

— — —

DT NN VBD NNS IN DT DT NN CC DT NN .  
The pound extended losses against both the dollar and the euro .

[ 1, 1, 0, ...]  
word-2=pound word-1=extended word+1=hello

# Encode as a bag of features

— — —

DT NN VBD NNS IN DT DT NN CC DT NN .

The pound extended losses against both the dollar and the euro .

“losses”/NNS:


{“left-bigram”: “pound extended”,

“right-bigram”: “against both”,

“context”: “extended losses against both”}

# Encode as a bag of features

---




“eat”/VBP:

[	1,	0,	0,	1,	0,	1,	0,	0,	...]
	nsubj	acl	relcl	dobj	pobj	prep	punct	xcomp	

# Encode as a bag of features

---




“eat”/VBP:

- *nsubj\_We, dobj\_pizza, prep\_with*
- *nsubj\_PRP, dobj\_NN, prep\_IN*
- *nsubj\_We, dobj\_pizza, prep\_with\_pobj\_anchovy*

# Encode as a bag of features

“fleas”/NNS:

{“label”: “NP”,  
“anc-left”: “PP”,  
“anc-right”: “S”,  
“span-width”:1}





# Example

<https://bit.ly/2KNsiLJ>

or

[https://github.com/mariana-scorp/esscass-2019-nlp/blob/master/  
2-features/feature-encoding.ipynb](https://github.com/mariana-scorp/esscass-2019-nlp/blob/master/2-features/feature-encoding.ipynb)

## 4. A few words about ngrams (optional)

# What are ngrams

---

Ngram - a contiguous sequence of  $n$  items from a given text.

# What are ngrams

---

Ngram - a contiguous sequence of  $n$  items from a given text.

So, if  $n = 3$ :

<S> Why did n't you listen to me ? </S>




# What are ngrams

---

Ngram - a contiguous sequence of  $n$  items from a given text.

So, if  $n = 3$ :

<S> Why did n't you listen to me ? </S>



# What are ngrams

---

Ngram - a contiguous sequence of  $n$  items from a given text.

So, if  $n = 3$ :

<S> Why did n't you listen to me ? </S>

# What are ngrams

---

Ngram - a contiguous sequence of  $n$  items from a given text.

So, if  $n = 3$ :

<S> Why did n't you listen to me ? </S>

# What are ngrams

---

Ngram - a contiguous sequence of  $n$  items from a given text.

So, if  $n = 3$ :

<S> Why did n't you listen to me ? </S>



# What are ngrams

---

Ngram - a contiguous sequence of  $n$  items from a given text.

So, if  $n = 3$ :

<S> Why did n't you listen to me ? </S>


# What are ngrams

---

Ngram - a contiguous sequence of  $n$  items from a given text.

So, if  $n = 3$ :

<S> Why did n't you listen to me ? </S>




# What are ngrams

---

Ngram - a contiguous sequence of  $n$  items from a given text.

So, if  $n = 3$ :

<S> Why did n't you listen to me ? </S>



# Token ngrams

---

Usually  $1 \leq n \leq 5$ .

<S> Why did n't you listen to me ? </S>

**n = 1:** (<S>), (Why), (did), (n't), (you), (listen), (to), (me), (?)...

**n = 2:** (<S> Why), (Why did), (did n't), (n't you), (you listen), (listen to)...

**n = 3:** (<S> Why did), (Why did n't), (did n't you), (you listen to)...

...

# Character Ngrams

---

<S> Why did n't you listen to me ? </S>

For words:

$n = 3$ : (<w> W h), (W h y), (h y </w>), (<w> d i), (d i d), (i d n), (d n ')...

For sentences:

$n = 3$ : (W h y), (h y \_), (y \_ d), (\_ d i), (d i d), (i d n), (d n '), (n ' t)...

# POS Ngrams

---

<S> Why did n't you listen to me ? </S>

<S> WDT VBD RB PRP VB TO PRP . </S>

POS:

**n = 3:** (<S>, WDT, VBD), (WDT, VBD, RB), (VBD, RB, PRP), (RB, PRP, VB)...

Token+POS:

**n = 2:** (<S>\_<S>, Why\_WDT), (Why\_WDT, did\_VBD), (did\_VBD, n't\_RB)...

Token or POS:

**n = 3:** (<S>, WDT, did), (WDT, did, RB), (did, RB, PRP), (RB, PRP, listen)...

# Tree Ngrams

---

Head+dependency:

*listen\_nsubj*


*listen\_nsubj\_you*

*listen\_prep\_to\_pobj\_me*

Head+POS+dependency:

*listen/VB\_nsubj*

*listen/VB\_nsubj\_you/PRP*



# Ngrams usage

- Speech recognition
- Text generation
- Autocompletion



google autocomplete is|

google autocomplete is **funny**

google autocomplete is **not working**

google autocomplete is **not working in firefox**

google autocomplete is **annoying**

google autocomplete is **slow**

google autocomplete islam

google autocomplete isn't working



# Ngrams usage

- Speech recognition
- Text generation
- Autocompletion
- Handwriting recognition
- Spelling correction
- (and GEC in general)

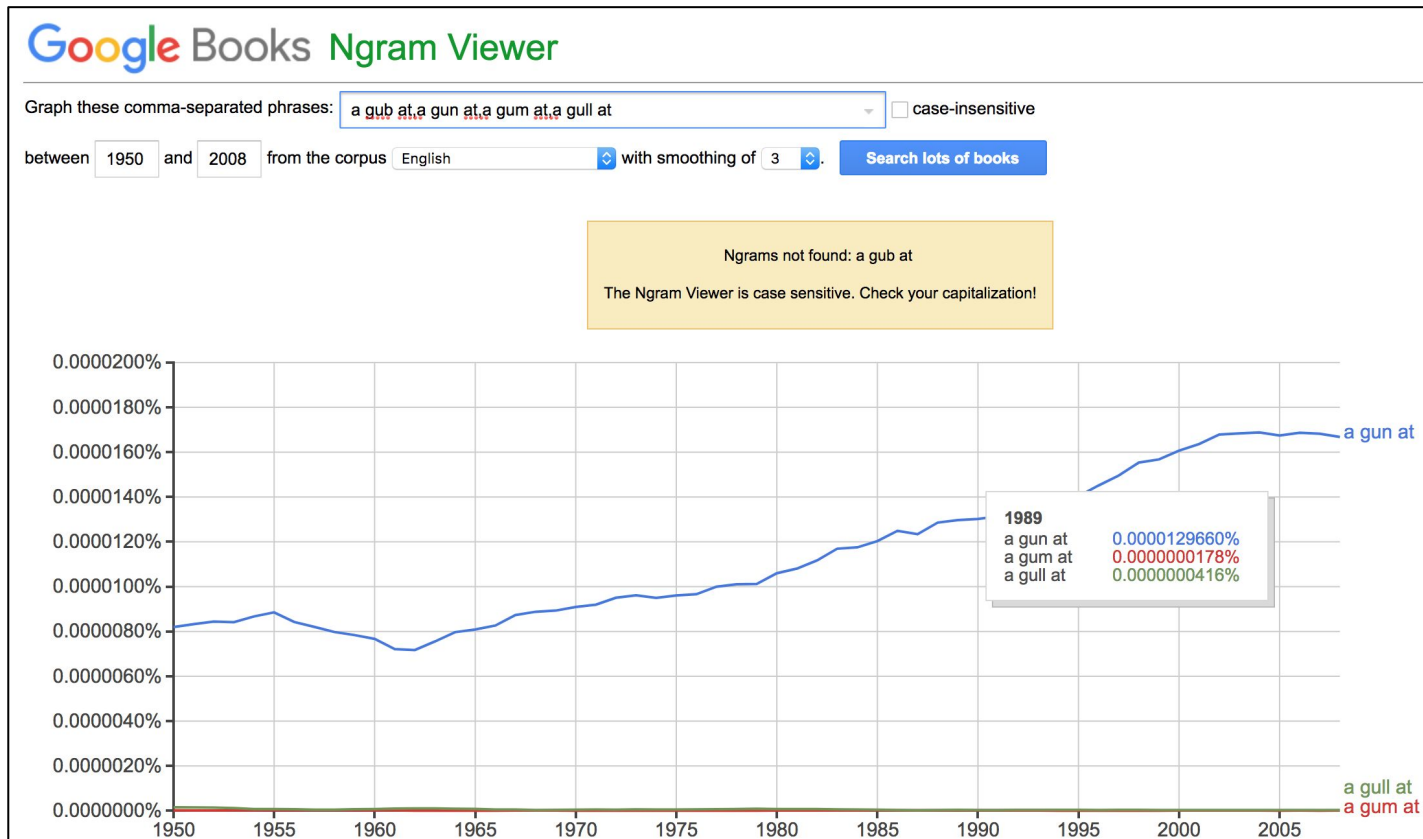


# Ngrams as a feature

---

Frequency or  
probability:

a gub at  
a gun at  
a gum at  
a gull at



# Ngrams as a feature

---

Frequency or probability

at	0.1
by	0.2
for	0.1
He will take our place <b>in</b> the line. →	<b>0.3</b>
from	0.0
to	0.1
with	0.1

# Ngrams as a feature


---

Conditional probability

$$P(w_n|w_{n-1}) = \frac{C(w_{n-1}w_n)}{C(w_{n-1})}$$

# Where to get ngrams

- [1 mln of 2/3/4/5-ngrams from COCA](#) for free
- [Google ngrams](#) (and [how to download](#))
- [Google syntactic ngrams](#)
- collect on your own



## 5. Logistic Regression (optional)

# Logistic Regression

---

Logistic regression - a discriminative linear model used for binary classification.

- *like Perceptron*, it's linear
- *like Naive Bayes*, it extracts a set of weighted features, takes logs, and combines them linearly
- *unlike Naive Bayes*, it's discriminative

$$z = \left( \sum_{i=1}^n w_i x_i \right) .$$


# Logistic Regression

---

A  $[0; 1]$  function would be handy:  $y = 1$  if  $p(y=1|x) > 0.5$ .

Sigmoid function:


$$\begin{aligned} P(y=1) &= \sigma(w \cdot x + b) \\ &= \frac{1}{1 + e^{-(w \cdot x + b)}} \end{aligned}$$






# Logistic Regression: multiclass

one vs. rest



multinomial (MaxEnt)



# Logistic Regression

---

For multinomial logistic regression, use softmax:

$$p(c|x) = \frac{\exp\left(\sum_{i=1}^N w_i f_i(c, x)\right)}{\sum_{c' \in C} \exp\left(\sum_{i=1}^N w_i f_i(c', x)\right)}$$

# Logistic Regression: example

---

Welcome to St . Paul 's Cathedral !



[Is this period a sentence end?]

# Logistic Regression: example

---

Welcome to St . Paul 's Cathedral !

 [Is this period a sentence end?]

$y$ : {is-end, is-not-end}

$x$ : {"word+1\_is\_cap", "word-1=kittens", "word-1=St", "tag-1=PRP", "tag+1=JJ"}

# Logistic Regression: example

---

Welcome to St . Paul 's Cathedral !



[Is this period a sentence end?]

$y$ : {is-end, is-not-end}


$x$ : {"word+1\_is\_cap", "word-1=kittens", "word-1=St", "tag-1=PRP", "tag+1=JJ"}

$x_j$ : [1, 0, 1, 0, 0]

# Logistic Regression: example

---

Welcome to St . Paul 's Cathedral !

 [Is this period a sentence end?]

$y$ : {is-end, is-not-end}

$x$ : {"word+1\_is\_cap", "word-1=kittens", "word-1=St", "tag-1=PRP", "tag+1=JJ"}


$x_j$ : [1, 0, 1, 0, 0]

$w_{\text{is-end}}$ : [2.9, 2.5, -0.9, 0, 0]

$w_{\text{is-not-end}}$ : [0.5, -0.7, 2.9, 0, 0]

# Logistic Regression: example

---  
Welcome to St . Paul 's Cathedral !

 [Is this period a sentence end?]

$y$ : {is-end, is-not-end}

$x$ : {"word+1\_is\_cap", "word-1=kittens", "word-1=St", "tag-1=PRP", "tag+1=JJ"}

$x_j$ : [1, 0, 1, 0, 0]

$w_{\text{is-end}}$ : [2.9, 2.5, -0.9, 0, 0]

$w_{\text{is-not-end}}$ : [0.5, -0.7, 2.9, 0, 0]

$$P(\text{is-end}|x_j) = e^{2.9-0.9} / (e^{2.9-0.9} + e^{0.5+2.9}) = 0.2$$

$$P(\text{is-not-end}|x_j) = e^{0.5+2.9} / (e^{2.9-0.9} + e^{0.5+2.9}) = 0.8$$

# Logistic Regression: weights

---

Learn weights:

- start with a vector of zeros
- move towards the gradient
- to maximize the probability / minimize the loss function

$$\hat{w} = \underset{w}{\operatorname{argmax}} \sum_j \log P(y^{(j)} | x^{(j)})$$



## 6. Error correction use case

# Error correction use case


<https://bit.ly/31GFZmd>

or

[https://github.com/mariana-scorp/esscass-2019-nlp/blob/master/  
2-features/adjective-vs-adverb.ipynb](https://github.com/mariana-scorp/esscass-2019-nlp/blob/master/2-features/adjective-vs-adverb.ipynb)



Thank you !



Any questions ?

[mariana.romanyshyn@grammarly.com](mailto:mariana.romanyshyn@grammarly.com)